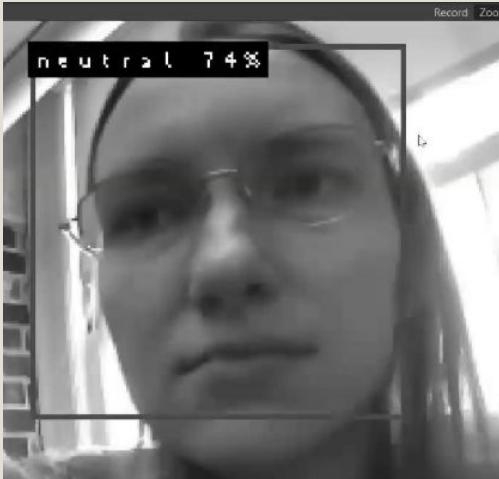


# TinyVision

---

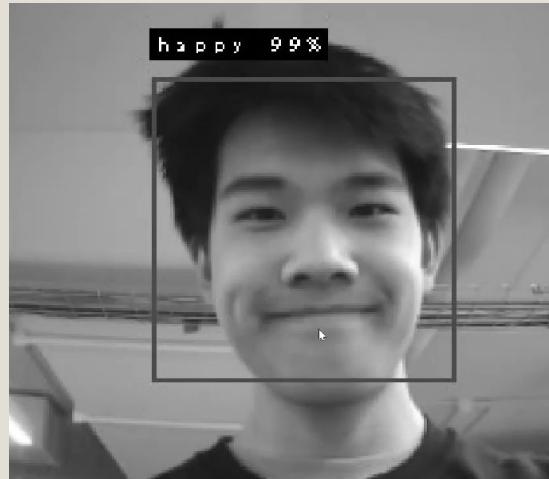
Masha, Rohina, Jordan

# Team & Roles



**Maria**

- Hardware setup and integration
- Model training



**Jordan**

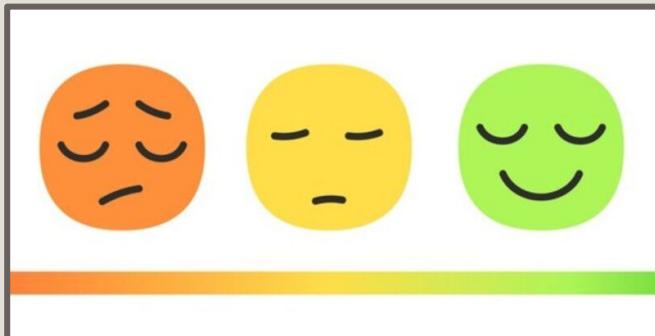
- Model training
- Hardware & demo testing
- Model architecture research



**Rohina**

- Model training
- Dataset research and prep

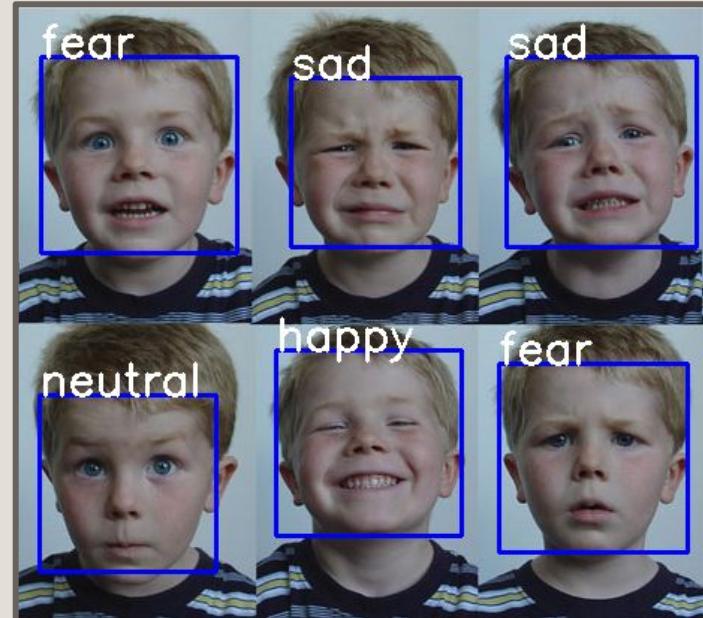
# Agenda



1. Introduction to FER
2. FER13 & CNN Discussion
3. RAF-DB & Tradeoffs
4. FER-Plus & Optimizations
5. Final Results
6. Conclusions

# Facial Expression Recognition (FER)

- Develop **real-time** system to detect **facial emotions**
- Positive vs. negative or emotion classes
  - Neutral
  - Happy
  - Angry
  - Sad
  - Surprise
  - ...
- Applications in **healthcare** to improve patient care
  - Therapy
  - Physical examinations
  - Research studies
- Need performance and energy **efficiency** + low **latency**



<https://sefiks.com/2018/01/10/real-time-facial-expression-recognition-on-streaming-data/>

# Goals

- Original goal: **> 80% accuracy** in real-time detection of **3+ emotion classes** at **>20 FPS**
- Made FPS goal more realistic after research:
  - OpenMV Cam H7 Plus “running time ... using [MobileNetV2] is around 0.3 seconds on average” with custom database (Asmara et al., 2024)
  - New latency goal of **<300 ms or >3 FPS**



# Research: datasets, existing models, & achievability

- Reviewed FER-2013 and CK Extended datasets
- Studied CNN & MobileNet-based FER models
- Prior work shows FER runs on OpenMV H7
- Confirms real-time edge deployment is feasible

# Hardware: OpenMV Cam H7 Plus

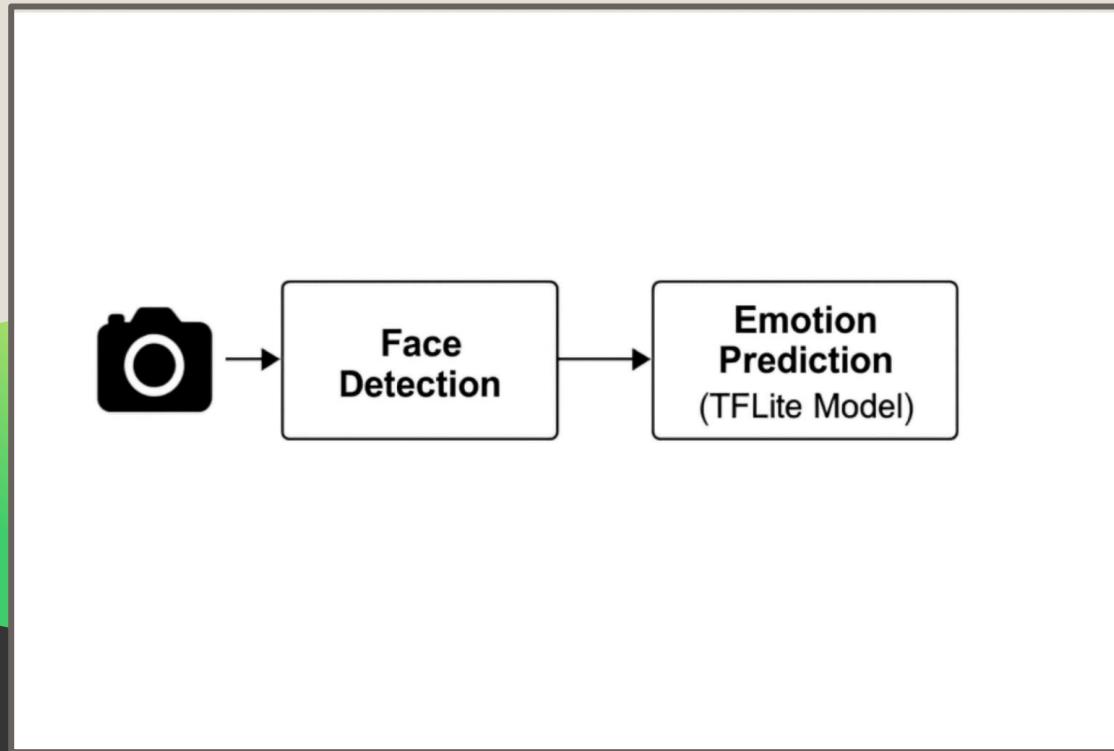
- ARM Cortex-M7 microcontroller
- On-board camera for real-time vision
- Supports TensorFlow Lite models
- Designed for low-power edge AI



# Training Considerations

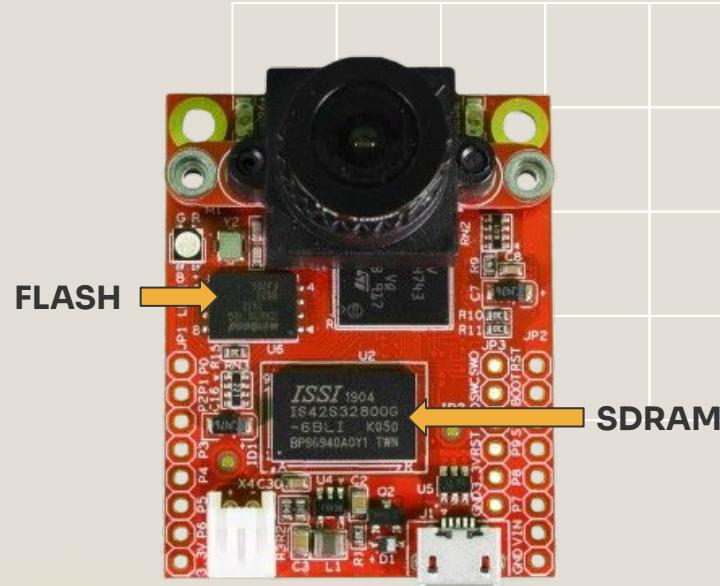
- Initial models achieved moderate accuracy
- Accuracy improved with fewer emotion classes
- Public datasets limit real-world generalization
- Lighting and image quality strongly affect accuracy

# FER Model Pipeline

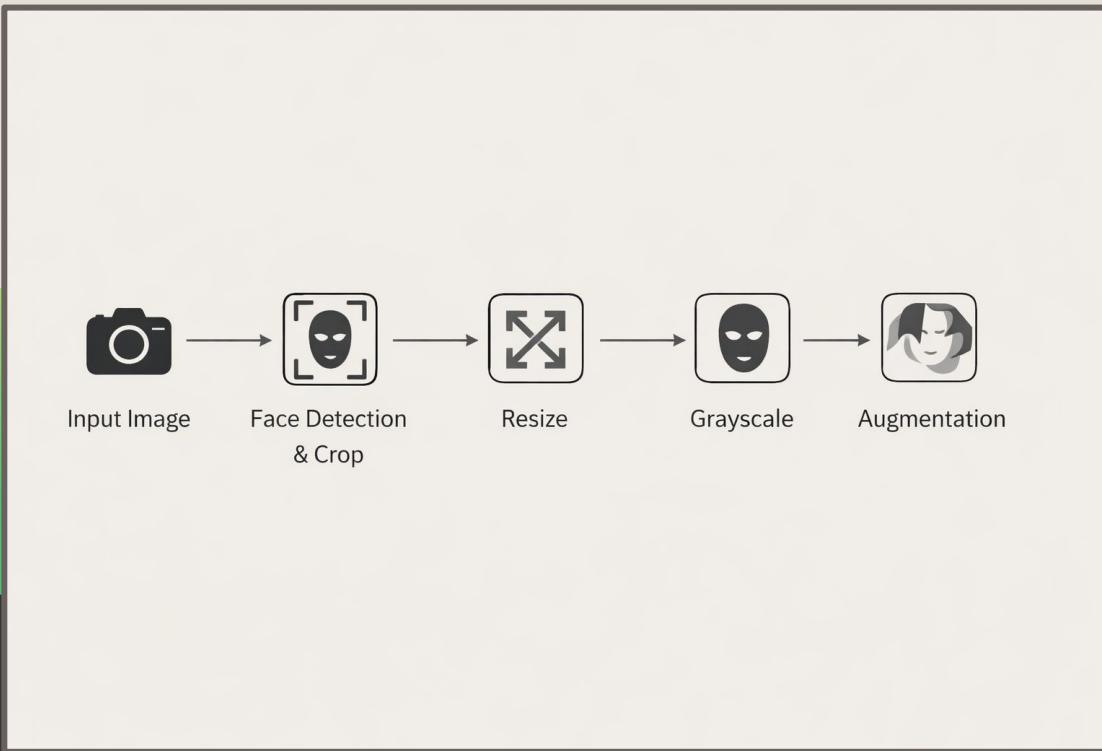


# Constraints

- Image size
  - 48 x 48 or 96 x 96
  - Tradeoff acceptable for emotion detection task
- Grayscale
- Minimizing TFLite arena w/ shallower network layers
- 32 MB extern. Flash
- <16 MB from 32 MB peak RAM available
  - Frame buffers
  - TensorArena → params, ops
  - Overhead

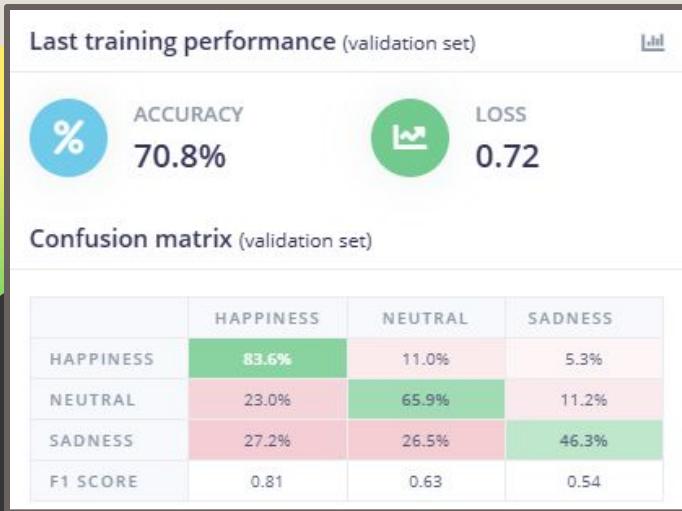


# Preprocessing Steps



# MobileNetV2

- Transfer learning with the lowest possible alpha value of 0.35
- Pretrained on ImageNet
- First trained with Edge Impulse, which limits us to 1 hour of compute, or around 25-30 epochs...



ACCURACY  
56.21%

Metrics for Transfer learning

METRIC	VALUE
Area under ROC Curve ⓘ	0.87
Weighted average Precision ⓘ	0.71
Weighted average Recall ⓘ	0.72
Weighted average F1 score ⓘ	0.71

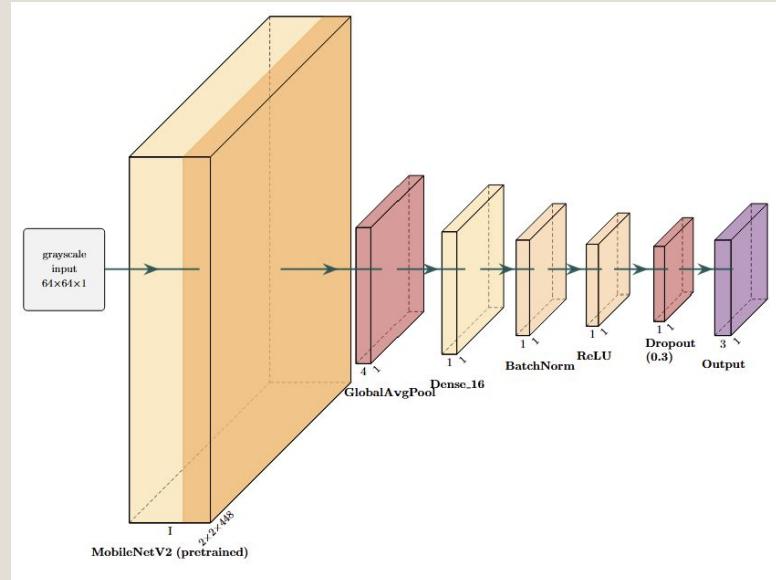
Confusion matrix

	HAPPINESS	NEUTRAL	SADNESS	UNCERTAIN
HAPPINESS	74.4%	5.8%	1.4%	18.4%
NEUTRAL	14.4%	43.4%	3.4%	38.8%
SADNESS	14.4%	10.5%	29.3%	45.8%
F1 SCORE	0.79	0.54	0.43	

# Can We Do... Better?

- Only around **4 MB of 32 MB** available to load in the model
- Attempting to improve model via more epochs, additional layers, QAT, or larger images/enhancement
- Ran many trials, many did not fit....
- Trained model locally, then used Edge Impulse EON compiler to quantize

Memory before loading model:  
Free memory: 4330176 bytes (4228.69 KB)  
Allocated memory: 7104 bytes (6.94 KB)



## Results

Model version: ⓘ Quantized (int8) ▾

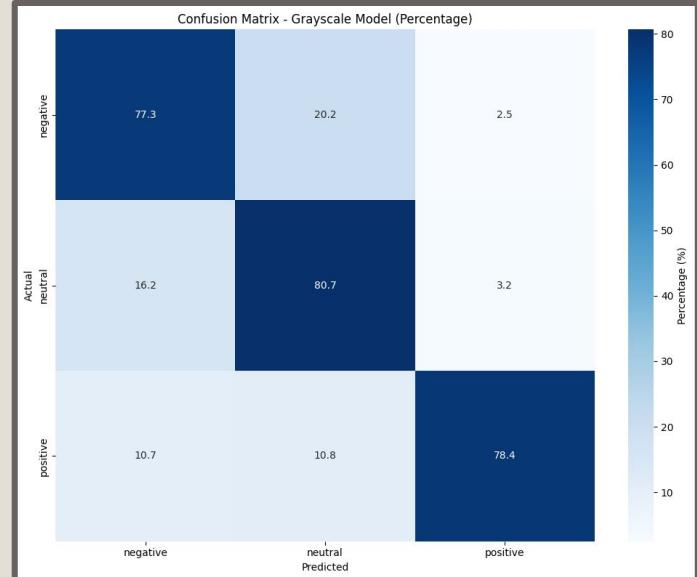
 ACCURACY  
67.95%

### Metrics for Pretrained learn

METRIC	VALUE
Area under ROC Curve ⓘ	0.93
Weighted average Precision ⓘ	0.80
Weighted average Recall ⓘ	0.77
Weighted average F1 score ⓘ	0.77

### Confusion matrix

	POSITIVE	NEGATIVE	NEUTRAL	UNCERTAIN
POSITIVE	47.6%	23.4%	4.5%	24.5%
NEGATIVE	4.1%	74.1%	4.7%	17.1%
NEUTRAL	2.7%	3.1%	84.6%	9.6%
F1 SCORE	0.62	0.69	0.88	



Validation Accuracy: 0.7869  
F1 Score: 0.7903  
Precision: 0.8015  
Recall: 0.7870

The screenshot shows the OpenMV Cam H7 Plus IDE interface. On the left, the code editor displays Python code for face detection, including imports for OpenCV and the OpenMV library, and a loop that processes frames and prints the elapsed time and FPS. The central part of the interface is the 'Frame Buffer' viewer, which shows a grayscale video frame of a person's face. A bounding box highlights the detected face area, and the word 'neutral' is displayed above it along with a confidence percentage of '74%'. Below the frame buffer, there are two tabs: 'Search Results' and 'Serial Terminal'. The Serial Terminal tab is active, showing the real-time output of the script, which consists of repeated lines of text indicating the processing of each frame. At the bottom of the screen, system status information is provided, including the board (OpenMV Cam H7 Plus), sensor (OV5640), firmware version (4.7.0 - [latest]), serial port (COM3), drive (D:/), and FPS (14.7).

# MobileNetV2 transfer learning demo

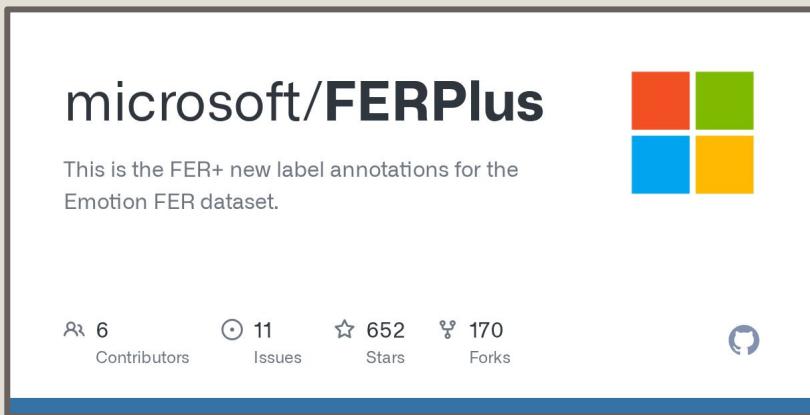
# MobileNetV2 transfer learning demo

The image shows the OpenMV Cam software interface. On the left, a code editor displays Python code for a face detection and emotion recognition application. The code includes imports for `time` and `faces`, initializes a clock, and loads a model. It then enters a loop where it prints the elapsed time and FPS, and classifies the detected faces as 'positive', 'neutral', or 'neutral'. On the right, a video frame shows a person's face being processed by the algorithm. At the bottom, a status bar provides hardware information: Board: OpenMV Cam H7 Plus, Sensor: OV5640, Firmware Version: 4.7.0 - [ latest ], Serial Port: COM3, Drive: D:/, and FPS: 13.3.

```
File Edit Tools Window Help
untitled_2.py x |
  faces
18
19 clock = time.clock()
20
21 # =====
22 # LOAD MODEL AND LABELS
Serial Terminal | ↴ ⌂ ⌂
Elapsed t: 46 ms; FPS: 21.73913
positive
Elapsed t: 46 ms; FPS: 21.73913
neutral
Elapsed t: 47 ms; FPS: 21.2766
neutral
Elapsed t: 46 ms; FPS: 21.73913
neutral
Elapsed t: 47 ms; FPS: 21.2766
neutral
Elapsed t: 47 ms; FPS: 21.2766
positive
Elapsed t: 46 ms; FPS: 21.73913
neutral
Elapsed t: 47 ms; FPS: 21.2766
neutral
X Search Results Serial Terminal
Board: OpenMV Cam H7 Plus Sensor: OV5640 Firmware Version: 4.7.0 - [ latest ] Serial Port: COM3 Drive: D:/ FPS: 13.3
```

MobileNetV2 transfer learning demo

# FERPlus & Optimizations



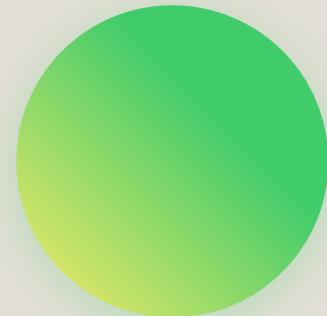
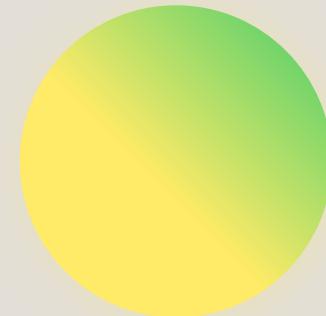
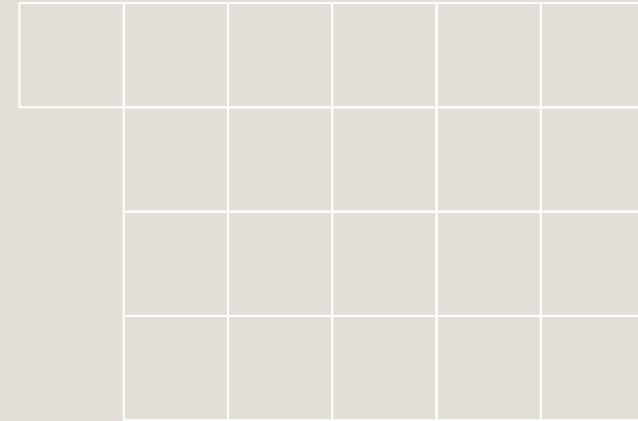
Higher accuracy achievable using new label annotations.  
<https://github.com/microsoft/FERPlus>

## FERPlus: Better Labeling

- New, less noisy, crowd-sourced labels (2016)
- Probability distribution among 8 classes (new contempt class)
- Removed bad samples

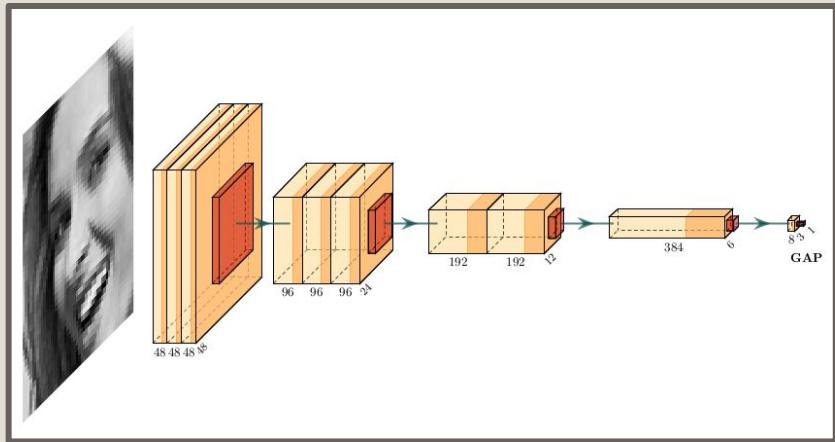
1	Usage	Image name	neutral	happiness	surprise	sadness	anger	disgust	fear	contempt	unknown	NF
2	Training	fer000000.png	4	0	0	1	3	2	0	0	0	0
3	Training	fer000001.png	6	0	1	1	0	0	0	0	0	2
4	Training	fer000002.png	5	0	0	3	1	0	0	0	0	1
5	Training	fer000003.png	4	0	0	4	1	0	0	0	0	1
6	Training	fer000004.png	9	0	0	1	0	0	0	0	0	0
7	Training	fer000005.png	6	0	0	1	0	0	0	1	1	1
8	Training	fer000006.png	2	0	0	8	0	0	0	0	0	0
9	Training	fer000007.png	0	10	0	0	0	0	0	0	0	0
10	Training	fer000008.png	0	10	0	0	0	0	0	0	0	0
11	Training	fer000009.png	0	0	6	0	0	0	4	0	0	0
12	Training	fer000010.png	2	0	0	0	8	0	0	0	0	0
13	Training	fer000011.png	10	0	0	0	0	0	0	0	0	0
14	Training	fer000012.png	5	0	0	3	0	0	0	0	0	2
15	Training	fer000013.png	9	0	0	1	0	0	0	0	0	0
16	Training	fer000014.png	0	10	0	0	0	0	0	0	0	0
17	Training	fer000015.png	0	0	6	0	1	0	3	0	0	0
18	Training	fer000016.png	4	6	0	0	0	0	0	0	0	0
19	Training		0	0	0	0	0	0	0	0	0	10
20	Training	fer000018.png	1	0	2	4	2	0	0	0	1	0
21	Training	fer000019.png	6	1	0	0	3	0	0	0	0	0

Higher accuracy achievable using new label annotations. <https://arxiv.org/abs/1608.01041>



# Testing a Promising FER Model

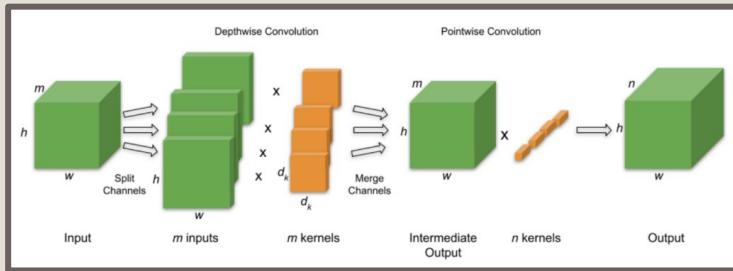
- Adam Wiącek's trained **87%** model is **~174k params, 244 KB size**
  - **334 ms** inference on Cam H7 Plus
- SeparableConv2D → BatchNorm and/or SpatialDropout →  
LeakyReLU → MaxPooling2D



Continuation of Adam Wiącek's master's thesis on "Deep learning in facial emotion recognition".  
<https://github.com/vicksam/fer-model>

Layer (type)	Output Shape	Param #
separable_conv2d_20	Separab (None, 48, 48, 48)	185
batch_normalization_18	(Bactc (None, 48, 48, 48)	192
leaky_re_lu_18 (LeakyReLU)	(None, 48, 48, 48)	0
separable_conv2d_21	Separab (None, 48, 48, 48)	2784
batch_normalization_19	(Bactc (None, 48, 48, 48)	192
leaky_re_lu_19 (LeakyReLU)	(None, 48, 48, 48)	0
separable_conv2d_22	Separab (None, 48, 48, 48)	2784
batch_normalization_20	(Bactc (None, 48, 48, 48)	192
max_pooling2d_8	(MaxPooling2 (None, 24, 24, 48)	0
spatial_dropout2d_8	(Spatial (None, 24, 24, 48)	0
leaky_re_lu_20 (LeakyReLU)	(None, 24, 24, 48)	0
separable_conv2d_23	Separab (None, 24, 24, 96)	5136
batch_normalization_21	(Bactc (None, 24, 24, 96)	384
leaky_re_lu_21 (LeakyReLU)	(None, 24, 24, 96)	0
separable_conv2d_24	Separab (None, 24, 24, 96)	10176
batch_normalization_22	(Bactc (None, 24, 24, 96)	384
leaky_re_lu_22 (LeakyReLU)	(None, 24, 24, 96)	0
separable_conv2d_25	Separab (None, 24, 24, 96)	10176
batch_normalization_23	(Bactc (None, 24, 24, 96)	384
max_pooling2d_9	(MaxPooling2 (None, 12, 12, 96)	0
spatial_dropout2d_9	(Spatial (None, 12, 12, 96)	0
leaky_re_lu_23 (LeakyReLU)	(None, 12, 12, 96)	0
separable_conv2d_26	Separab (None, 12, 12, 192)	19488
batch_normalization_24	(Bactc (None, 12, 12, 192)	768
leaky_re_lu_24 (LeakyReLU)	(None, 12, 12, 192)	0
separable_conv2d_27	Separab (None, 12, 12, 192)	38784
batch_normalization_25	(Bactc (None, 12, 12, 192)	768
max_pooling2d_10	(MaxPooling (None, 6, 6, 192)	0
spatial_dropout2d_10	(Spatial (None, 6, 6, 192)	0
leaky_re_lu_25 (LeakyReLU)	(None, 6, 6, 192)	0
separable_conv2d_28	Separab (None, 6, 6, 384)	75840
batch_normalization_26	(Bactc (None, 6, 6, 384)	1536
max_pooling2d_11	(MaxPooling (None, 3, 3, 384)	0
spatial_dropout2d_11	(Spatial (None, 3, 3, 384)	0
leaky_re_lu_26 (LeakyReLU)	(None, 3, 3, 384)	0
separable_conv2d_29	Separab (None, 3, 3, 8)	3464
global_average_pooling2d_2	( (None, 8)	0
<hr/>		
Total params: 173,537		
Trainable params: 171,137		
Non-trainable params: 2,400		

# Optimizing for Cam H7 Plus



Depthwise Separable Convolution filters applied to input to produce output. Assume stride=1 and padding= $d=2$ . Source: [Efficient Deep Learning Book](#)



Heavy bias towards the first 5 emotions in Adam Więcek's FER model. <https://github.com/vicksam/fer-model>

Reduce params & compute time:

**SepableConv2D** replaces Conv2D

- Depthwise spatial conv on each channel + pointwise mixing of channels

**SpatialDropout2D** replaces regular Dropout

- Drop entire 2D feature maps (not point-elements) → promote fmap independence

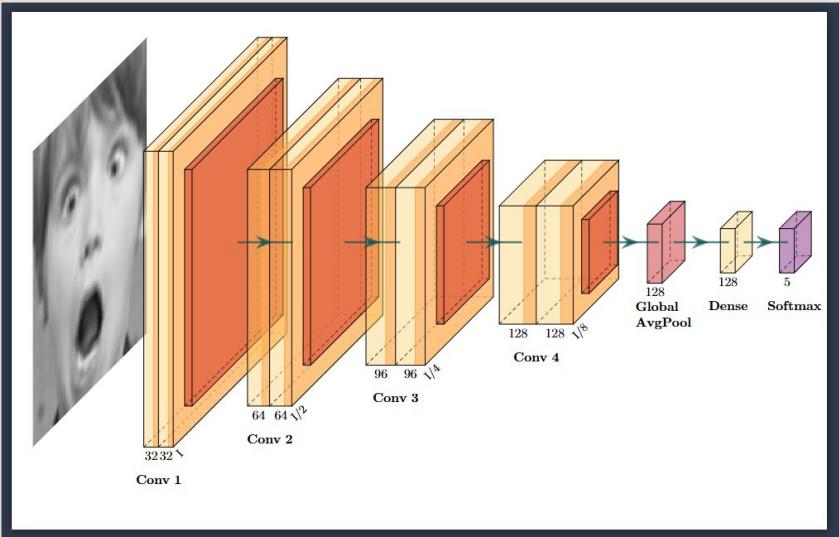
**GlobalAveragePooling2D** replaces Flatten

- Averages + pools + reshapes

- 
- **Reduced to 5 classes:** neutral, happiness, surprise, sadness, angriness

# Retraining Using Custom CNNs

- 1st layer = regular Conv2D, all others **SeparableConv2D**
- **Regular ReLU** replaces LeakyReLU → faster convergence
- **Added Dense** layer after GAP → increased learning performance



Our optimized CNN model based on Adam Wiącek's architecture  
<https://github.com/vicksam/fer-model>

Layer (type)	Output Shape	Param #
conv1_1 (Conv2D)	(None, 48, 48, 32)	320
bn1_1 (BatchNormalization)	(None, 48, 48, 32)	128
sepconv1_2 (SeparableConv2D)	(None, 48, 48, 32)	1,344
bn1_2 (BatchNormalization)	(None, 48, 48, 32)	128
pool1 (MaxPooling2D)	(None, 24, 24, 32)	0
spatial_dropout1 (SpatialDropout2D)	(None, 24, 24, 32)	0
sepconv2_1 (SeparableConv2D)	(None, 24, 24, 64)	2,480
bn2_1 (BatchNormalization)	(None, 24, 24, 64)	256
sepconv2_2 (SeparableConv2D)	(None, 24, 24, 64)	4,736
bn2_2 (BatchNormalization)	(None, 24, 24, 64)	256
pool2 (MaxPooling2D)	(None, 12, 12, 64)	0
spatial_dropout2 (SpatialDropout2D)	(None, 12, 12, 64)	0
sepconv3_1 (SeparableConv2D)	(None, 12, 12, 96)	6,816
bn3_1 (BatchNormalization)	(None, 12, 12, 96)	384
sepconv3_2 (SeparableConv2D)	(None, 12, 12, 96)	10,176
bn3_2 (BatchNormalization)	(None, 12, 12, 96)	384
pool3 (MaxPooling2D)	(None, 6, 6, 96)	0
spatial_dropout3 (SpatialDropout2D)	(None, 6, 6, 96)	0
sepconv4_1 (SeparableConv2D)	(None, 6, 6, 128)	13,280
bn4_1 (BatchNormalization)	(None, 6, 6, 128)	512
sepconv4_2 (SeparableConv2D)	(None, 6, 6, 128)	17,664
bn4_2 (BatchNormalization)	(None, 6, 6, 128)	512
pool4 (MaxPooling2D)	(None, 3, 3, 128)	0
spatial_dropout4 (SpatialDropout2D)	(None, 3, 3, 128)	0
global_avg_pool (GlobalAveragePooling2D)	(None, 128)	0
dropout_gap (Dropout)	(None, 128)	0
fc1 (Dense)	(None, 128)	16,512
dropout_fc (Dropout)	(None, 128)	0
predictions (Dense)	(None, 7)	903

Total params: 76,711 (299.65 KB)

Trainable params: 75,431 (294.65 KB)

Non-trainable params: 1,280 (5.00 KB)

## Training graphs



## Results

ACCURACY  
73.56%

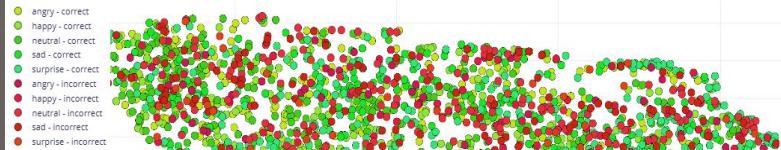
### Metrics for Classifier

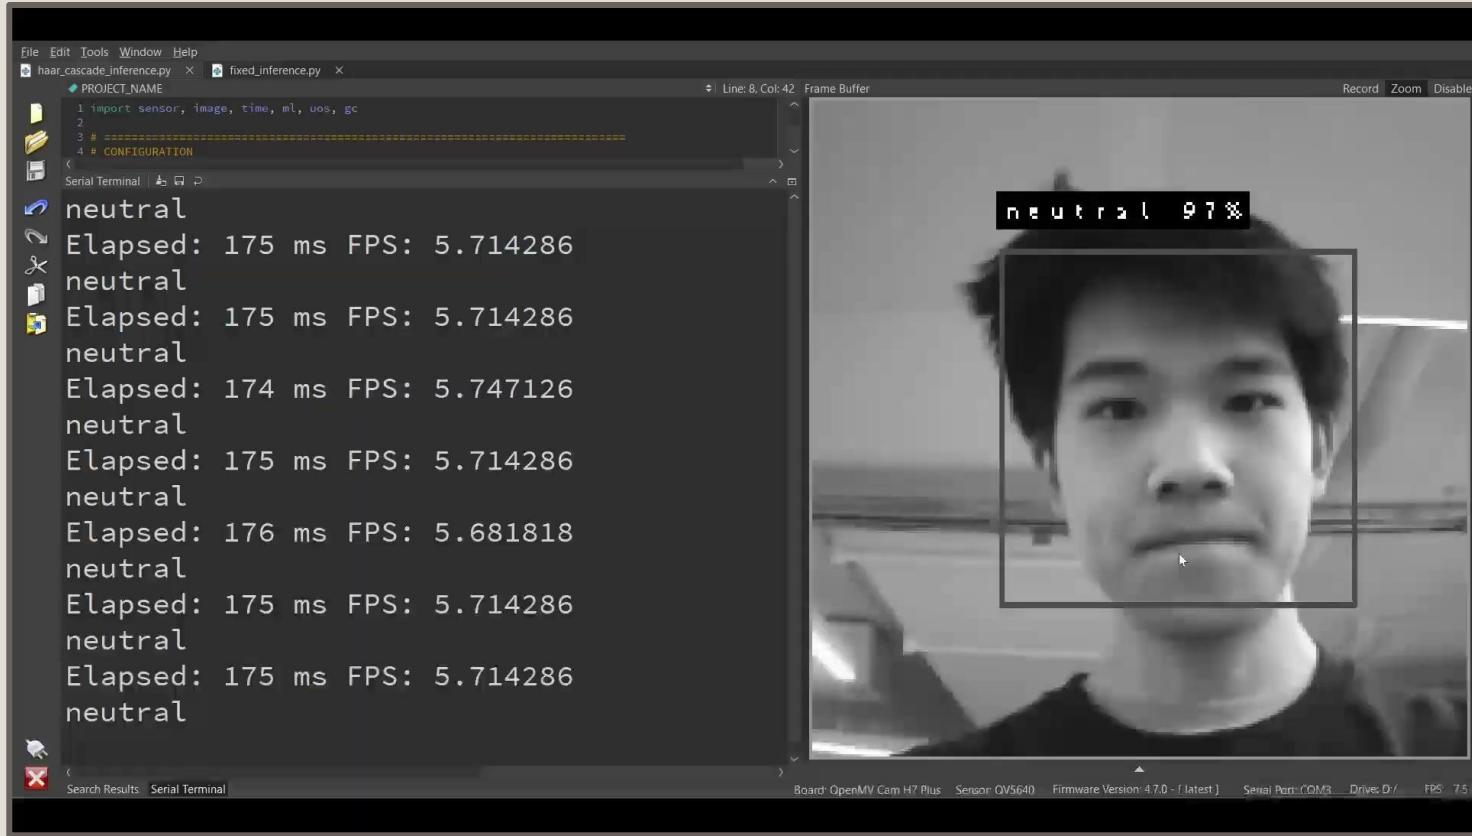
METRIC	VALUE
Area under ROC Curve ⓘ	0.96
Weighted average Precision ⓘ	0.79
Weighted average Recall ⓘ	0.79
Weighted average F1 score ⓘ	0.79

### Confusion matrix

	ANGRY	HAPPY	NEUTRAL	SAD	SURPRISE	UNCERTAIN
ANGRY	79.8%	1.0%	3.6%	2.5%	2.3%	10.9%
HAPPY	3.4%	85.8%	3.2%	1.7%	3.0%	7.9%
NEUTRAL	4.3%	2.6%	62.7%	8.2%	4.2%	16.5%
SAD	6.1%	1.4%	12.4%	59.4%	1.6%	19.1%
SURPRISE	4.2%	0.7%	1.7%	0.7%	86.6%	6.0%
F1 SCORE	0.80	0.86	0.72	0.67	0.87	

### Feature explorer ⓘ





Custom CNN demo; Peak RAM usage: 158.5 KB; Flash usage 160.6KB



Custom CNN demo; Peak RAM usage: 158.5 KB; Flash usage 160.6KB

# Future Work

- Achieved goals of **80%+ training acc** and **>3 FPS inference time** in real-world testing
- **Standardize** real-world testing to calculate meaningful “accuracy”
  - Control lighting, contrast, background
  - Run tests across different people
- Explore **hyperparameter optimization**; mainly tweaked by observations in this project
  - Learning rate
  - Regularization strength
  - Maximum arch. depth
  - Kernel type / size
  - Batch size



# References, Thank You!

OpenMV Cam H7 Plus Documentation:

<https://openmv.io/products/openmv-cam-h7-plus>

<https://docs.openmv.io/library/omv.image.html#class-haarcascade-feature-descriptor>

Facial Expression Recognition Literature:

<https://github.com/vicksam/fer-model>

<https://github.com/microsoft/FERPlus>

Asmara et al., 2024; <https://joiv.org/index.php/joiv/article/view/2299/0>

Other: <https://www.tensorflow.org/>

